

Who Should Be in the Loop?

Allocating Expert Review for Qualitative Coding with Language Models

Anonymous ACL submission

Abstract

Large language models can assist qualitative coding, but specialist review does not scale easily. We present **WarrantRoute**, a source-grounded protocol for testing which reviewer expertise is needed for particular qualitative judgments and for allocating review under a fixed expert-time budget. The planned study will compare researchers without specialist expertise, qualitative methods experts, and domain experts on natural outputs and controlled variants targeting unsupported evidence, source concentration, counterevidence loss, contextual flattening, and unsupported abstraction. It will evaluate whether a policy frozen before held-out testing can route items without hiding less frequent or negative cases, and whether selected feedback repairs defects without introducing new ones. The current manuscript reports the protocol and a measurement-pipeline qualification on entirely fictional packets; it reports no real-corpus model, rater, routing, or repair result.

Manuscript skeleton: gray fields and blank table cells must be replaced from the frozen analysis export.

1 Introduction

Large language models (LLMs) can suggest codes, organize codebooks, and apply qualitative categories across large collections of text (Dai et al., 2023; Wang et al., 2026; Matveyenko et al., 2026). Researchers often review these outputs, edit suggestions, explain disagreements, and guide later analysis. Most evaluations still treat their judgments as one general source of human feedback, even when the task calls for different kinds of knowledge.

A qualitative methods expert may recognize that a topic label has been presented as an interpretive theme. A domain expert may notice that an analysis has lost a cultural distinction or omitted a clinically important exception. A researcher without either specialization may still identify an unsupported

quotation or an overly broad claim. A single average score conceals these differences and provides little guidance about how to use limited expert time.

Evaluations of LLM supported clinical thematic analysis remain fragmented across expert review and automatic similarity measures (Yi et al., 2025a). The usual measures do not resolve this problem. More than one code space may be defensible (Chen et al., 2026), agreement between humans and LLMs may not predict blind expert preference (Liu et al., 2026), and model confidence can be poorly calibrated (Li et al., 2026). Audits have also found missing participant evidence, summaries presented as quotations, and failures that depend on local context (Parkington et al., 2025; Bang et al., 2026; Subbiah et al., 2026). Provenance identifies the source of a passage. It does not show that the passage supports the interpretation or represents the range of views in the corpus.

The planned experiment treats review as an *allocation problem*: who should examine a qualitative interpretation and its supporting evidence? **WarrantRoute** will collect judgments from researchers without specialist expertise, qualitative methods experts, and domain experts. The study will include natural model outputs and controlled variants with known defects. A simple policy will send each item to no specialist, a qualitative methods expert, a domain expert, or both. The policy will use information available before specialist review, and all comparisons will receive the same amount of expert time. We will freeze the policy before testing it on a new domain and model family. Figure 1 summarizes the framework.

The planned study is designed to contribute:

1. a source-grounded benchmark that isolates five validity problems without assuming a single correct thematic interpretation;
2. a human-evaluation design that adapts evidential credibility and plural-reference cover-




Figure 1: Overview of the planned WarrantRoute evaluation. An interpretable policy will route source-grounded theme and evidence packets to researchers alone, a qualitative methods expert, a domain expert, or both within a fixed time budget. The study will evaluate routing on Cache2 (project alias: CaChE) and Dreaddit and test whether expert feedback repairs errors without creating new ones.

age from prior work and adds study-specific, source-grounded operationalizations of voice-and-boundary preservation and scope calibration; and

3. a planned held-out domain test that compares several allocation rules at equal cost and measures whether selected feedback repairs a defect without creating a new one.

The present manuscript reports the protocol and a synthetic measurement-pipeline qualification, not a real-corpus human-evaluation result.

2 Related Work

Human feedback in qualitative analysis.

Human-in-the-loop systems let analysts revise codes and rationales (Dai et al., 2023), pair-code and discuss disagreements (Gao et al., 2024), convert repeated feedback into persistent instructions (Wang et al., 2026), and retain positionality and revision histories (Ye et al., 2026). Multi-agent systems extend this workflow: Auto-TA automates source-to-theme generation (Yi et al., 2025c), SFT-TA uses supervised fine-tuned agents (Yi et al., 2025b), and TAMA incorporates cardiac expert guidance (Xu et al., 2026). These studies establish scalable collaboration, but their evaluations emphasize alignment, usability, control, or aggregate output quality rather than reviewer specialization.

Review allocation and evaluator dependence.

Selective-review methods route items using confi-

dence and ensemble diversity (Zhao and Liu, 2025) or rationale disagreement (Tajik et al., 2026), although confidence can be miscalibrated or anchor-sensitive (Li et al., 2026). Evaluator identity also changes conclusions: blind expert verification can diverge from agreement metrics (Liu et al., 2026), and specialists who read source material can differ from nonspecialists who see isolated themes (Moskalewicz et al., 2026). Adjacent biomedical work evaluates generated protocol pseudocode against a generated reference (Yi et al., 2024), but not the source warrant of qualitative interpretations. Prior work therefore studies allocation signals and evaluator perspective separately, not which expert role a source-linked packet requires under matched expert time.

Source grounding and artifact quality.

Provenance systems preserve links among quotations, codes, themes, and revisions (Ye et al., 2026; Yi et al., 2026), yet models can omit human-identified excerpts, substitute summaries for quotations, or produce paraphrased, truncated, and hallucinated evidence (Parkington et al., 2025; Bang et al., 2026). Computational thematic-analysis evaluations have operationalized credibility through human acceptance and semantic alignment and, later, binary LLM judgments (Qiao et al., 2025; Wang et al., 2026); OpenCodes instead measures coverage relative to a pooled code space (Chen et al., 2026). These constructs are complementary: cited evidence may support a claim that omits dissent, while broad coverage does not ensure appropri-

ate scope. WarrantRoute therefore adapts evidential credibility and plural-reference coverage, adds packet-level operationalizations of voice-and-boundary preservation and scope calibration, and reports all four separately. The open comparison is whether methods and domain expertise detect complementary source-grounded failures, and whether expertise-aware routing improves detection over expertise-agnostic allocation at the same cost.

3 Task and Study Design

3.1 Evidence packets and controlled variants

The planned unit of analysis is a **theme and evidence packet**. Each future real-data packet will contain a qualitative claim, a short explanation, and exact excerpts only from an access-controlled, pseudonymized analysis copy after the required two-person privacy and permissions review, with project-pseudonymous speaker or document identifiers, local context, and a summary of source coverage. Raters will judge whether the evidence supports the claim and preserves important differences, counterevidence, and source voice. The task is a bounded source-warrant and representation assessment. It will not assess a complete reflexive thematic analysis or assume that one thematic account is universally correct (Braun and Clarke, 2021a,b).

The study will sample natural LLM outputs and construct variants vetted by experts. Each variant will change one property frozen in the preregistration:

1. **unsupported evidence**: a topically related excerpt does not warrant the claim;
2. **source concentration**: wording implies breadth while evidence comes disproportionately from one participant or document;
3. **counterevidence loss**: a rare, negative, or boundary case is removed;
4. **contextual flattening**: a consequential cultural or domain specific qualification is erased; or
5. **unsupported abstraction**: a local observation is generalized beyond its evidence.

An independent construction panel will check the original packet, the intended defect, and the absence of a second manipulation. Pilot items will not enter the main analysis, and each rater will see at most one version of a base packet.

3.2 Artifact-quality measures

We will preregister four artifact-quality measures, separately from the routing and repair outcomes below. Two adapt measures used in prior computational thematic analysis. **Evidential credibility** asks whether the core interpretive meaning is directly supported by the cited evidence in local context. Unlike prior binary LLM judgments (Qiao et al., 2025; Wang et al., 2026), the planned primary measure is a blinded human rating. **Plural-reference coverage** is the percentage of concepts in a frozen reference space pooled from multiple independent analysts that are recovered by at least one generated code or theme (Chen et al., 2026). The pool will preserve alternative concepts rather than designate one analyst’s themes as truth, and blinded experts will adjudicate many-to-many conceptual matches.

We propose two study-specific, packet-level operationalizations. **Voice-and-boundary preservation** assesses whether consequential differences, dissent, counterevidence, less-frequent or boundary cases, and contextual qualifications survive synthesis. **Scope calibration** assesses whether claim breadth, strength, polarity, and causal language are proportionate to the supplied evidence. Raters will score these two measures and evidential credibility from 1 to 5. The analysis will retain the ordinal scores and report ratings of 4 or 5 as descriptively adequate; “cannot judge” will be reported separately. Plural-reference coverage will be calculated only on complete natural-output inventories, separately for codes and themes, and will not be treated as a 1–5 rating. We will interpret the four measures only after deterministic quotation, attribution, assignment, and provenance checks. Because coverage is recall-like, we will report generated-concept volume and human-audited unsupported or redundant concepts so that verbosity cannot masquerade as quality. We will not combine the four measures into a composite. Full anchors and matching rules appear in Appendix B.4.

3.3 Datasets and leakage controls

The project maintains four candidate corpus lanes that differ in domain, interaction structure, and provenance. Dreddit’s official training split is provisionally assigned to development, its official test split to a secondary in-domain audit, and the 11 Cache2 focus groups (project alias: CaChE) to the planned held-out cross-domain confirma-

tory evaluation. CaSiNo and the AMI scenario-meeting subset are the current source-only dialogue candidates, covering task-oriented dyadic negotiation and elicited multiparty meetings. Their exact source snapshots have been acquired and integrity checked, but neither lane yet has adapter-validated records, an approved model or rater pathway, or a frozen experimental role or split. Here, “held out” describes the planned procedure, not model-training provenance.

Table 1 reports each lane at its current preparation stage. Dreaddit and CaChe counts come from restricted processed copies, whereas CaSiNo and AMI counts are verified source inventories. The source only rows are not common experimental denominators and are not sampling eligible. Definitions, provenance, access conditions, and calculations appear in Appendix B.1.

Under the planned protocol, only Dreaddit training data will inform prompts, packet rules, controlled variants, routing features, and thresholds. We will freeze these choices before test generation, preserve complete source groups, quarantine exact training duplicates that overlap Dreaddit test text, and cluster CaChe analyses by focus group and transcript-local speaker. CaSiNo and AMI will not inform these choices while their experimental roles remain unassigned. Packets will use the same context limits across adopted evaluation datasets. We will retain prompts, decoding settings, run identifiers, source spans, and revisions. Protected text will be processed only under approved terms. The complete leakage and generation protocol appears in Appendix B.1.

3.4 Raters and annotation

The planned study will recruit three groups. **Qualitative methods experts** will have demonstrated experience conducting qualitative analysis. **Domain experts** will have research, professional, or relevant lived experience in the source domain. **Researchers without specialist expertise** will be able to apply the study rubric but will not meet either expert definition for the assigned item. Dual expertise will be recorded as a separate status. Item construction, initial rating, policy development, and final adjudication will be handled by different people wherever staffing permits.

All raters will complete the same tutorial and review balanced subsets in a randomized incomplete-block design. For each packet they will rate evidential credibility, voice-and-boundary preservation,

scope calibration, and confidence on scales from 1 to 5. The common quality anchors will be 1 (not adequate), 2 (major problems), 3 (partially adequate; material revision needed), 4 (mostly adequate; minor issues only), and 5 (fully adequate); confidence will use separate uncertainty anchors. They will also choose accept, revise, reject, or escalate; indicate whether they cannot judge the item; provide a short rationale and, if needed, request a particular kind of expertise. The system will record review time. “Cannot judge” will remain a separate response category, with no conversion to a midpoint or a quality failure. Qualitative methods experts and domain experts will submit independent ratings before adjudication. When their disagreement reflects defensible interpretations, the item will remain plural or ambiguous.

We will choose the final sample size after a separate pilot by simulating the crossed structure of raters and packets. The preregistration will state the number of unique raters, base packets, ratings per group, exclusions, and the smallest effect or noninferiority margin that the design can support.

4 Routing by Expertise

The planned routing method will use information available before specialist review: researchers’ ratings of evidential credibility, voice-and-boundary preservation, and scope calibration, plus disagreement, confidence, abstention, the requested kind of expertise, response time, and nonsensitive provenance measures such as excerpt count, source concentration, and quotation exactness. Plural-reference coverage will be evaluation-only because it relies on the frozen analyst reference space. The router will not use manipulation labels, expert scores from the test set, protected participant attributes, or test outcomes.

We will fit an interpretable regularized model that chooses among no escalation, review by a qualitative methods expert, review by a domain expert, or review by both. We will select its thresholds on the development data to maximize recall of serious errors within a fixed budget measured in expert minutes. We will then compare five conditions:

1. **researchers only**: no specialist review;
2. **both expert groups**: a qualitative methods expert and a domain expert review every item;
3. **random**: route the same budget at random;

Dataset	Domain and context	Language	Sources	Text units	Text count	Status
Dreaddit (Turcan and McKeown, 2019)	Social media stress; short texts	English	2,929 posts	3,553 source; 3,532 eligible	302,509 eligible words	Processed; review pending
Cache2 (CaChe) (Mehta, 2024)	Sexual and reproductive health; focus groups	Translated English	11 groups	3,944 retained; 3,076 eligible	98,735 eligible words	Processed; review pending
CaSiNo (Chawla et al., 2021)	Campsite negotiation; dyadic text dialogue	English	1,030 dialogues	11,919 ordinary messages	228,675 source words	Source only
AMI scenario (Carletta et al., 2006)	Product design; four party meetings	Spoken English	138 meetings	69,258 source segments	793,764 source <w> elements	Source only

Table 1: Dataset inventory at the current preparation stage. Dreaddit and CaChe eligible counts apply deterministic content rules but remain pending manual privacy review and project authorization. CaSiNo source counts use ordinary messages and whitespace delimited words and exclude 2,378 structured negotiation events. AMI counts upstream speech segment and orthographic word XML elements. Source only rows are neither processed nor sampling eligible, and the differing unit types are not common analysis denominators.

4. **uncertainty only**: route using confidence, disagreement, and abstention; and
5. **WarrantRoute**: predict both whether escalation is needed and which role is relevant.

Every condition will use the same observed expert time. The planned confirmatory outcomes on the CaChe cross-domain evaluation are recall of serious errors, false acceptance, and expert minutes per correctly escalated error; the same outcomes will be reported on the Dreaddit in-domain audit as secondary results. These policy-level outcomes are distinct from the four artifact-quality measures. A secondary experiment will use one fixed revision model and compare no feedback, feedback from researchers, feedback selected by WarrantRoute, and feedback from both expert groups. A separate blinded panel will judge whether the revision fixes the target defect and whether it creates a new unsupported claim, omission, or distortion.

5 Evaluation and Analysis

The planned analysis addresses three questions. First, for which errors do qualitative methods experts or domain experts differ materially from researchers without specialist expertise? Second, can the initial ratings identify those cases within a fixed budget of expert time? Third, does the selected feedback repair defects in a new domain and model family?

For the first question, separate cumulative-link mixed-effects models will analyze evidential credibility, voice-and-boundary preservation, and scope calibration; logistic mixed-effects models will analyze decisions and defect detection. Models will in-

clude packet version, evaluator group, error family, domain, and their preregistered interactions, with random effects for rater and base packet. Plural-reference coverage will be estimated separately at the complete-run level, not from sampled packets or controlled twins. For the second question, we will compare policies on the cross-domain evaluation with cluster-bootstrap intervals and a preregistered noninferiority contrast against review by both expert groups. For the third question, a logistic mixed model will relate successful repair to feedback condition and error family. Dreaddit results will be secondary; the CaChe analysis will become confirmatory only after its manifest and analysis plan are frozen and preregistered.

We will report effect sizes and uncertainty intervals. Ordinal Krippendorff’s α , Gwet’s AC2, raw agreement, and confidence calibration will describe measurement behavior; they will not establish validity on their own (McDonald et al., 2019). We will freeze and preregister confirmatory contrasts, multiplicity, and the CaChe packet manifest before generating model outputs for that corpus. Development and audit results, comparisons of natural and controlled items, and other robustness checks will appear in Appendix D. Table 7 is reserved for reliability, and Table 8 for transfer and routing-feature ablations. Confidence and router calibration will be reported separately in the appendix.

6 Results

This section separates the completed synthetic measurement-pipeline qualification from the planned real-corpus reporting shell. The qualification used four fictional templates from an ear-

lier corpus design, including KODIS-shaped and CANDOR-shaped proxies. It tests whether the four measures can be calculated reproducibly; it provides no evidence about Dreaddit, CaCHe, CaSiNo, AMI, or any other real corpus. Replace every later blank from a single frozen analysis export and reconcile all denominators before writing the final abstract. Report the CaCHe confirmatory results before the Dreaddit audit and development performance.

6.1 Synthetic measurement-pipeline qualification

We applied the paper-aligned measures to 60 outputs generated from entirely fictional packets: five recent model variants, four corpus-shaped packets, and three independent runs per model and packet. Three blinded model judges independently scored evidential credibility, voice-and-boundary preservation, and scope calibration using the manuscript anchors. For this engineering pilot, the benchmark’s frozen designed constructs instantiate the reference space; each judge matched every reference separately to generated codes and themes, and a two-of-three panel rule determined recovery. Table 2 reports all four measures.

The pilot exposes a ceiling effect that would be hidden by reporting adequacy alone. GPT-5.6 Sol, GPT-5.5, and GPT-5.6 Terra received 100% adequacy on all three ordinal measures, while their deterministic integrity pass rates were 10/12, 9/12, and 5/12. Code coverage ranged only from 93.7% to 95.2%, and theme coverage from 85.7% to 93.7%. We therefore retain the four constructs as the source-grounded artifact profile but do not use them alone for model selection: integrity failures are exclusionary, full ordinal distributions and controlled-defect sensitivity are reported, and coverage is accompanied by output-volume and unsupported-concept diagnostics. Under that gate-first rule, GPT-5.6 Sol remains the remediation candidate rather than a frozen paper model.

6.2 Study accounting

The final manuscript will report [N TOTAL RATERS] recruited raters and [N RATINGS] analyzed judgments over [N PACKETS] packets. [REPORT EXCLUSIONS, BALANCE, AND THE TEST SET]. Table 3 is reserved for study counts by experimental role.

6.3 Artifact-quality profile

The final manuscript will report CaCHe credibility, voice-and-boundary, and scope-calibration adequacy rates, plus code and theme plural-reference coverage, only after the confirmatory analysis is frozen and run. It will report these measures with the integrity prerequisite, denominators, and uncertainty, and provide full ordinal distributions, cannot-judge rates, matching uncertainty, unmatched reference concepts, and generated-concept volume in Appendix B.4.

6.4 Where expertise changes judgments

[SUMMARIZE THE INTERACTION BETWEEN EVALUATOR GROUP AND FAILURE FAMILY, INCLUDING UNCERTAINTY]. Table 4 reports predicted probabilities of detecting a defect. [STATE WHETHER THE TWO EXPERT GROUPS PROVIDED COMPLEMENTARY, REDUNDANT, OR INCONCLUSIVE JUDGMENTS].

6.5 Routing under matched expert time

After the expert-minute budget is preregistered and the study is run, the final manuscript will report [PRIMARY TEST RESULT], [EFFECT VERSUS UNCERTAINTY ROUTING WITH 95% CONFIDENCE INTERVAL], and [NONINFERIORITY RESULT VERSUS BOTH EXPERT GROUPS]. Table 5 is reserved for the five routing policies under the matched expert-time budget. Results must be reported by failure family if an average gain conceals misses on negative cases.

6.6 Does routed feedback repair defects?

The final manuscript will report [RATE AND CONFIDENCE INTERVAL] for packets repaired by WarrantRoute feedback without a new error, alongside [RESEARCHER FEEDBACK RATE] for researcher feedback and [BOTH EXPERT GROUPS RATE] for feedback from both expert groups. Table 6 is reserved for repair outcomes, new errors, credibility changes, and expert time. [REPORT PERFORMANCE RELATIVE TO EXPERT MINUTES AND THE MOST IMPORTANT FAILURE CASE].

7 Discussion

[STATE WHICH EVALUATOR GROUP ADDED VALUE FOR WHICH FAILURE TYPES, WITH EFFECTS]. [STATE

Model	Integrity pass	Evidential adequacy	Voice/boundary adequacy	Scope adequacy	Code coverage	Theme coverage	Serious error
GPT-5.6 Sol	10/12	100.0 [100.0, 100.0]	100.0 [100.0, 100.0]	100.0 [100.0, 100.0]	93.7 [89.1, 98.4]	85.7 [78.5, 93.4]	0.0 [0.0, 0.0]
GPT-5.4	9/12	100.0 [100.0, 100.0]	97.2 [91.7, 100.0]	100.0 [100.0, 100.0]	93.7 [89.1, 98.4]	93.7 [89.1, 98.4]	5.6 [0.0, 13.9]
GPT-5.5	9/12	100.0 [100.0, 100.0]	100.0 [100.0, 100.0]	100.0 [100.0, 100.0]	93.7 [89.1, 98.4]	92.1 [87.3, 96.8]	0.0 [0.0, 0.0]
GPT-5.6 Luna	8/12	80.6 [69.4, 91.7]	80.6 [69.4, 91.7]	88.9 [80.6, 97.2]	95.2 [90.9, 100.0]	85.7 [79.7, 91.9]	8.3 [0.0, 16.7]
GPT-5.6 Terra	5/12	100.0 [100.0, 100.0]	100.0 [100.0, 100.0]	100.0 [100.0, 100.0]	95.2 [90.9, 100.0]	88.9 [84.1, 93.7]	0.0 [0.0, 0.0]

Table 2: Paper-aligned measurement pilot on entirely fictional packets. Ordinal columns report the percentage of 36 blinded model-judge ratings per model scored 4 or 5. Code and theme coverage report panel-majority recovery across 63 frozen reference opportunities per level and model. Brackets are 95% packet-by-run cluster-bootstrap intervals. Integrity is deterministic. All three judges and candidates are from one provider; these are measurement-pipeline results, not real-corpus or human-evaluation results.

Quantity	Dev.	In domain	Cross domain
Source participants/documents	—	—	—
Base packets / controlled twins	—	—	—
Researchers without specialist expertise	—	—	—
Qualitative methods experts	—	—	—
Domain experts	—	—	—
Independent ratings	—	—	—
Median review minutes	—	—	—

Table 3: Planned reporting shell for study counts by experimental role. Raters and packets will be reported as unique counts.

WHETHER THE TWO EXPERT GROUPS PROVIDED COMPLEMENTARY JUDGMENTS]. [REPORT THE PRACTICAL COST AND EXPERT-REVIEW COVERAGE CONSEQUENCE AT THE OBSERVED REVIEW TIMES]. Any resulting conclusions must be limited to the evidence-review task, failure families, and evaluator groups actually studied.

On the new domain and model family, **[STATE THE TRANSFER RESULT AND ITS INTERPRETATION]. [STATE WHAT THE RESULT IMPLIES FOR ROUTING BEYOND GENERIC UNCERTAINTY].** A useful review queue should preserve the reason for escalation and allow a reviewer to request another form of expertise. Persistent disagreement remains visible for deliberation.

8 Limitations

Any resulting estimates will reflect the recruited reviewers, whose judgments may reproduce disciplinary, cultural, or institutional assumptions. Controlled variants will isolate individual failure mechanisms, but naturally occurring errors may be subtler or combine several defects. Repeated judgments from the same people may increase precision without increasing the diversity of expertise. Routing can also create false reassurance around items that are not escalated, so the design includes

a random audit and failure-family-specific false-acceptance reporting. The task covers bounded evidence packets. A complete qualitative analysis also requires familiarization, memoing, reflexivity, and accountable interpretation. Generalization will be limited to English Reddit stress segments and English translations of focus groups from western Kenya, as well as **[MODEL FAMILIES AND RECRUITMENT SAMPLE]**. If the source-only dialogue lanes are later adopted, CaSiNo cannot support cross-dialogue participant counts or participant-disjoint splits from its public release, while AMI scenario meetings are elicited role play and transcript-only processing omits prosody, timing, overlap, and multimodal evidence.

9 Ethics Statement

Real text model processing and rater exposure begin only after a documented institutional determination and all source, model provider, rater access, retention, and release gates are satisfied; source collection approvals do not transfer. We treat all corpora as potentially identifying. Analysis copies are minimized and restricted, protected text goes only to approved services, and every excerpt shown to raters or released must pass a documented two-person privacy and permissions review. Raters participate voluntarily outside supervisory or evaluative relationships, with identifiers stored separately. Public artifacts contain only aggregate, synthetic, or cleared material. Routing supports research review, not clinical or participant decisions. Details appear in Appendix C.

10 Conclusion

WarrantRoute is designed to measure how reviewer expertise affects the detection and repair of unsupported qualitative interpretations under a fixed review budget. The current manuscript establishes

Failure family	Researchers	Methods experts	Domain experts	Largest contrast
Unsupported evidence	_____	_____	_____	_____
Source concentration	_____	_____	_____	_____
Counterevidence loss	_____	_____	_____	_____
Contextual flattening	_____	_____	_____	_____
Unsupported abstraction	_____	_____	_____	_____

Table 4: Predicted probability of detecting each failure, with 95% confidence intervals, by evaluator group.

Policy	Expert minutes	Serious error recall	False acceptance	Difference from both experts	Errors per expert hour
Researchers only	0	_____	_____	_____	Not applicable
Random	_____	_____	_____	_____	_____
Uncertainty routing	_____	_____	_____	_____	_____
WarrantRoute	_____	_____	_____	_____	_____
Both expert groups	_____	_____	_____	Reference	_____

Table 5: Planned reporting shell for routing performance on the CaChe confirmatory evaluation at matched budgets measured in expert minutes. Intervals will use a bootstrap clustered by base packet.

Feedback	Repair	New error	Credibility change	Expert minutes
None	_____	_____	_____	0
Researchers	_____	_____	_____	0
WarrantRoute	_____	_____	_____	_____
Both expert groups	_____	_____	_____	_____

Table 6: Planned reporting shell for blinded revision. A repair will count as successful only if it fixes the targeted defect without introducing an unsupported claim.

the protocol and reports only the fictional-packet measurement qualification; [STATE THE PRIMARY REAL-CORPUS RESULT AND ITS PRACTICAL IMPLICATION ONLY AFTER THE FROZEN STUDY IS COMPLETE].

References

- AMI Corpus. n.d. AMI corpus ethics and consent forms. <https://groups.inf.ed.ac.uk/ami/corpus/ethicsandconsent.shtml>. Official corpus documentation; accessed 25 August 2026.
- Enid Awiti, Sophie Young, Garazi Zulaika, Fredrick Odhiambo Otieno, Elizabeth Nyothach, Penelope A. Phillips-Howard, Supriya D. Mehta, and Linda Mason. 2024. “whenever i help her, i am also expecting her vagina in return”: A qualitative analysis to explore men’s and adolescent girls’ perceptions of the impact of the COVID-19 pandemic on the sexual behaviour and health of adolescent girls in rural western kenya. *BMJ Public Health*, 2(2):e001214. Peer-reviewed journal article; open access under CC BY.
- So Hyeon Bang, Soojeong Han, Meghan Reading Turchio, Melani Ellison, Stacey Dai, Mary Beth Happ, David Russell, and Ruth Masterson Creber. 2026. Developing and evaluating human-led and large language model-assisted hybrid deductive-inductive workflows for qualitative analysis. *Journal of the American Medical Informatics Association*. Peer-reviewed journal article.
- Virginia Braun and Victoria Clarke. 2021a. Can i use TA? should i use TA? should i not use TA? comparing reflexive thematic analysis and other pattern-based qualitative analytic approaches. *Counselling and Psychotherapy Research*, 21:37–47. Peer-reviewed journal article.
- Virginia Braun and Victoria Clarke. 2021b. One size fits all? what counts as quality practice in (reflexive) thematic analysis? *Qualitative Research in Psychology*, 18(3):328–352. Peer-reviewed journal article.
- Jean Carletta, Simone Ashby, Sébastien Bourban, Mike Flynn, Maël Guillemot, Thomas Hain, Jaroslav Kadlec, Vasilis Karaiskos, Wessel Kraaij, Melissa Kronenthal, Guillaume Lathoud, Mike Lincoln, Agnes Lisowska, Iain McCowan, Wilfried Post, Dennis Reidsma, and Pierre Wellner. 2006. The AMI meeting corpus: A pre-announcement. In *Machine Learning for Multimodal Interaction*, volume 3869 of *Lecture Notes in Computer Science*, pages 28–39. Springer.
- Kushal Chawla, Jaysa Ramirez, Rene Clever, Gale Lucas, Jonathan May, and Jonathan Gratch. 2021. CaSiNo: A corpus of campsite negotiation dialogues for automatic negotiation systems. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 3167–3185. Association for Computational Linguistics.
- John Chen, Alexandros Nikolaos Lotsos, Sihan Cheng, Lexie Zhao, Yanjia Zhang, Jessica Hullman, Bruce Sherin, Uri Wilensky, and Michael Horn. 2026. A computational method for measuring open codes in qualitative analysis. In *Findings of the Association for Computational Linguistics: ACL 2026*, pages 41740–41758. Peer-reviewed archival Findings paper.
- Shih-Chieh Dai, Aiping Xiong, and Lun-Wei Ku. 2023. LLM-in-the-loop: Leveraging large language model for thematic analysis. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 9993–10001. Peer-reviewed archival Findings paper.
- Jie Gao, Yuchen Guo, Gionnieve Lim, Tianqin Zhang, Zheng Zhang, Toby Jia-Jun Li, and Simon Tangi Perreault. 2024. CollabCoder: A lower-barrier, rigorous workflow for inductive collaborative qualitative analysis with large language models. In *Proceedings of the 2024 CHI Conference on Human Factors in Computing Systems*. Peer-reviewed CHI full paper.
- Hongming Li, Huan Kuang, and Anthony F. Botelho. 2026. Can we trust AI’s self-assessment? evaluating and improving LLM confidence calibration in educational dialogue coding. In *Artificial Intelligence in Education*. AIED 2026 conference work; formal proceedings carry 2027 metadata.
- Alex Liu, Lief Esbenshade, Michael Xiao, Victor Tian, Zachary Zhang, Kevin He, and Min Sun. 2026. Agreement is not quality: Blind expert verification of human and LLM qualitative coding when human consensus is not ground truth. *arXiv preprint arXiv:2607.28890*. Preprint; no peer-reviewed archival version verified as of 24 August 2026.
- Joseph Matveyenko, James Liu, John David Parsons, Ryan A. Brown, Alina I. Palimaru, Vipul Gupta, and Prateek Puri. 2026. Development and benchmarking of a blended human-AI qualitative research assistant. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics: Industry Track*. Peer-reviewed ACL Industry Track paper.
- Nora McDonald, Sarita Schoenebeck, and Andrea Forte. 2019. Reliability and inter-rater reliability in qualitative research: Norms and guidelines for CSCW and HCI practice. *Proceedings of the ACM on Human-Computer Interaction*, 3(CSCW):72:1–72:23. Peer-reviewed journal/proceedings article.
- Supriya Mehta. 2024. Cache2 focus group discussion transcripts AGYW and community males. University of Illinois Chicago Figshare. Dataset, version 1, CC BY 4.0.
- Marcin Moskalewicz, Anna Sterna, Karolina Drozd, Kacper Dudzic, Marek Pokropski, and Paula Flores. 2026. Computational phenomenology of borderline personality disorder: A comparative evaluation of LLM-simulated expert personas and human clinical experts. *arXiv preprint arXiv:2508.19008*. Workshop presentation reported; archival peer-review status not independently verified by 24 August 2026.
- Karisa Parkington, Bazen G. Teferra, Marianne Rouleau-Tang, Argyrios Perivolaris, Alice Rueda, Adam Dubrowski, Bill Kapralos, Reza Samavi, Andrew Greenshaw, Yanbo Zhang, Bo Cao, Yuqi Wu, Sirisha Rambhatla, Sridhar Krishnan, and Venkat Bhat. 2025.

687	Human vs. LLM-based thematic analysis for digital mental health research: Proof-of-concept comparative study. <i>arXiv preprint arXiv:2507.08002</i> . Preprint; no peer-reviewed archival version verified as of 24 August 2026.	745
688		746
689		747
690		748
691		749
692	Tingrui Qiao, Caroline Walker, Chris Cunningham, and Yun Sing Koh. 2025. Thematic-LM: A LLM-based multi-agent system for large-scale thematic analysis. In <i>Proceedings of the ACM on Web Conference 2025</i> , pages 649–658. ACM. Peer-reviewed conference paper.	750
693		751
694		752
695		753
696		
697		
698	Melanie Subbiah, Haaris Mian, Nicholas Deas, Ananya Mayukha, Dan P. McAdams, and Kathleen McKeown. 2026. Whose story gets told? positionality and bias in LLM summaries of life narratives. <i>arXiv preprint arXiv:2604.20131</i> . Preprint; no peer-reviewed archival version verified as of 24 August 2026.	754
699		755
700		756
701		757
702		758
703		759
704		760
705	Elham Tajik, Conrad Borchers, Bahar Shahrokhian, Sebastian Simon, Ali Keramati, Sonika Pal, and Sreecharan Sankaranarayanan. 2026. Disagreement as data: Reasoning trace analytics in multi-agent systems. In <i>Proceedings of the 16th International Learning Analytics and Knowledge Conference</i> . Peer-reviewed LAK paper.	761
706		
707		
708		
709		
710		
711		
712	Elsbeth Turcan and Kathleen McKeown. 2019. Dreddit: A reddit dataset for stress analysis in social media. In <i>Proceedings of the Tenth International Workshop on Health Text Mining and Information Analysis</i> , pages 97–107. Peer-reviewed workshop paper; the paper reports no ethics determination or dataset license.	762
713		763
714		764
715		765
716		766
717		767
718		768
719	Lei Wang, Min Huang, and Eduard Dragut. 2026. CentaurTA: A self-improving human-agents collaboration framework for thematic analysis. In <i>Findings of the Association for Computational Linguistics: ACL 2026</i> , pages 15871–15884. Peer-reviewed archival Findings paper.	769
720		770
721		
722		
723		
724		
725	Huimin Xu, Seungjun Yi, Terence Lim, Jiawei Xu, Andrew Well, Carlos Mery, Aidong Zhang, Yuji Zhang, Heng Ji, Keshav Pingali, Yan Leng, and Ying Ding. 2026. TAMA: A human-AI collaborative thematic analysis framework using multi-agent LLMs for clinical interviews. <i>ACM Transactions on Computing for Healthcare</i> . Peer-reviewed journal article.	771
726		772
727		773
728		774
729		775
730		776
731		777
732	Runlong Ye, Oliver Huang, Patrick Yung Kang Lee, Michael Liut, Carolina Nobre, and Ha-Kyung Kong. 2026. Reflexis: Supporting reflexivity and rigor in collaborative qualitative analysis through design for deliberation. In <i>Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems</i> . Peer-reviewed CHI full paper.	778
733		779
734		
735		
736		
737		
738		
739	Seungjun Yi, Jaeyoung Lim, and Juyong Yoon. 2024. ProtoMed-LLM: An automatic evaluation framework for large language models in medical protocol formulation. <i>arXiv preprint arXiv:2410.04601</i> . Preprint; presented at CHITA 2025; no peer-reviewed archival version verified as of 24 August 2026.	780
740		781
741		782
742		783
743		784
744		
	Seungjun Yi, Joakim Nguyen, Terence Lim, Andrew Well, Joseph Skrovan, Mehak Beri, YongGeon Lee, Kavita Radhakrishnan, Leqi Liu, Mia Markey, and Ying Ding. 2025a. Position: Thematic analysis of unstructured clinical transcripts with large language models. <i>arXiv preprint arXiv:2509.14597</i> . Preprint; presented at GenAI4Health at NeurIPS 2025; no peer-reviewed archival version verified as of 24 August 2026.	785
		786
		787
		788
		789
	Seungjun Yi, Joakim Nguyen, Huimin Xu, Terence Lim, Joseph Skrovan, Mehak Beri, Hitakshi Modi, Andrew Well, Leqi Liu, Mia Markey, and Ying Ding. 2025b. SFT-TA: Supervised fine-tuned agents in multi-agent LLMs for automated inductive thematic analysis. <i>arXiv preprint arXiv:2509.17167</i> . Preprint and ML4H 2025 manuscript; final archival status not verified as of 24 August 2026.	790
		791
		792
	Seungjun Yi, Joakim Nguyen, Huimin Xu, Terence Lim, Joseph Skrovan, Mehak Beri, Hitakshi Modi, Andrew Well, Carlos M. Mery, Yan Zhang, Mia K. Markey, and Ying Ding. 2026. Automated thematic analysis for clinical qualitative data: Iterative codebook refinement with full provenance. <i>arXiv preprint arXiv:2603.08989</i> . Preprint; manuscript reports submission to AMIA 2026; no peer-reviewed archival version verified as of 24 August 2026.	793
		794
		795
		796
		797
		798
	Seungjun Yi, Joakim Nguyen, Huimin Xu, Terence Lim, Andrew Well, Mia Markey, and Ying Ding. 2025c. Auto-TA: Towards scalable automated thematic analysis (TA) via multi-agent large language models with reinforcement learning. <i>arXiv preprint arXiv:2506.23998</i> . Preprint; presented as a non-archival paper at the ACL 2025 Student Research Workshop; no peer-reviewed archival version verified as of 24 August 2026.	799
		800
		801
		802
		803
		804
		805
		806
		807
		808
		809
		810
		811
		812
		813
		814
		815
		816
		817
		818
		819
		820
		821
		822
		823
		824
		825
		826
		827
		828
		829
		830
		831
		832
		833
		834
		835
		836
		837
		838
		839
		840
		841
		842
		843
		844
		845
		846
		847
		848
		849
		850
		851
		852
		853
		854
		855
		856
		857
		858
		859
		860
		861
		862
		863
		864
		865
		866
		867
		868
		869
		870
		871
		872
		873
		874
		875
		876
		877
		878
		879
		880
		881
		882
		883
		884
		885
		886
		887
		888
		889
		890
		891
		892
		893
		894
		895
		896
		897
		898
		899
		900
		901
		902
		903
		904
		905
		906
		907
		908
		909
		910
		911
		912
		913
		914
		915
		916
		917
		918
		919
		920
		921
		922
		923
		924
		925
		926
		927
		928
		929
		930
		931
		932
		933
		934
		935
		936
		937
		938
		939
		940
		941
		942
		943
		944
		945
		946
		947
		948
		949
		950
		951
		952
		953
		954
		955
		956
		957
		958
		959
		960
		961
		962
		963
		964
		965
		966
		967
		968
		969
		970
		971
		972
		973
		974
		975
		976
		977
		978
		979
		980
		981
		982
		983
		984
		985
		986
		987
		988
		989
		990
		991
		992
		993
		994
		995
		996
		997
		998
		999
		1000

and the AMI scenario-meeting subset are the two current source-only dialogue corpus candidates. Exact source snapshots for both have been acquired and integrity checked, but their experimental roles remain unassigned. Neither has an adapter-validated output, an approved model or rater pathway, a completed contextual privacy review, or a frozen project split.

Dreaddit is a publicly downloadable English Reddit corpus for stress detection. Its released labeled subset contains 3,553 human annotated segments from 2,929 posts across ten subreddits in five domains (Turcan and McKeown, 2019). The archive provides official train and test partitions and post identifiers, but no author identifier. It therefore supports segment to post provenance and document concentration checks, not participant counts or participant level claims. The author archive contains 2,838 train segments from 2,343 posts and 715 test segments from 586 posts. It contains no README or dataset license. We retain text, the stress label, subreddit, and sentence range in the restricted analysis copy. We omit original post identifiers, timestamps, engagement measures, the internally inconsistent confidence field, and unused lexical features. Project specific identifiers preserve post grouping without exposing the released post identifier. The official splits have no post identifier overlap, but three exact texts occur in both partitions under different identifiers and conflicting labels. We retain those rows for audit but make the three training copies and 18 later within train exact duplicates ineligible for sampling, leaving 2,817 development segments and 715 audit segments.

The Cache2 deposit, called CaChE in this project, contains 11 participant-numbered and selectively redacted focus-group transcripts about how COVID-19 affected schooling, sexual behaviour, and sexual and reproductive health in western Kenya (Mehta, 2024). The source study reports six groups with 54 adolescent girls and young women and five with 53 community men. Sessions used English, Swahili, Dholuo, or a mixture and were translated into English (Awiti et al., 2024). The deposit is version 1, DOI 10.25417/uic.26495884.v1, and CC BY 4.0. A fixed checksum allowlist verifies all 11 source files. Processing retains 3,575 participant turns and 369 collective turns; 3,076 participant turns meet the earlier three-word eligibility rule, but short and collective turns are no longer deleted and remain available as context. Each record has a stable focus-

group and turn identifier, a transcript-local speaker identifier, source line spans, source hashes, translation status, the preceding moderator question, and prior-turn references. Normalization yields 106 observed transcript-local participant keys, one fewer than the 107 participants reported by the source paper; we do not claim complete participant coverage until this is reconciled. Conflicting source dates are recorded but excluded from model and router features.

The acquired CaSiNo source is the official CC BY 4.0 repository snapshot pinned to commit 2f6ed4a. It contains 1,030 dyadic campsite negotiation dialogues, 11,919 ordinary messages, 2,378 structured negotiation events, and 228,675 whitespace delimited words in the ordinary messages (Chawla et al., 2021). The source paper reports 846 participants, but the public JSON provides only two dialogue local role tokens and no stable cross-dialogue participant token. We therefore reject the supplied 900/30/100 random split for WarrantRoute, do not reconstruct linkage from participant profiles or text, and leave the corpus unassigned unless trustworthy provider linkage becomes available or the entire corpus is preregistered in one unsplit role after the governance gates are satisfied. Any future adapter will retain ordinary dialogue text and turn lineage only. It will exclude structured events, annotations, demographics, personality measures, preferences, free text reasons, negotiation outcomes, and raw identifiers.

The acquired AMI manual annotation artifact contains a selectively extracted source inventory for 138 elicited scenario meetings in 35 meeting series. The four person meetings are English product design discussions. The selected files contain 793,764 orthographic word elements, 69,258 speech segment elements, and mappings for 140 global source speaker identifiers. These form 35 meeting and speaker connected components, 33 with four meetings and two with three meetings (Carletta et al., 2006). No audio, video, signals, slides, questionnaires, demographics, participant file, or nontranscript annotation category was extracted. The official filename labels the manual annotations as version 1.6.2, while the embedded README says release 1.7; we bind the source by its exact URL, filename, size, and checksum rather than asserting that those labels are equivalent. The corpus is not sampling eligible. A future transcript adapter must discard timing and raw source identifiers, and all 35 meeting and speaker components

must remain indivisible in any adopted split.

For Table 1, a source is one Reddit post, focus group, CaSiNo dialogue, or AMI meeting. Dreddit source units are released labeled segments, and its eligible units exclude three training copies that overlap test text plus 18 later within-train duplicates. CaChe retained units include participant and collective turns, while its eligible count applies the earlier three word participant-turn rule. CaSiNo units are ordinary source messages and exclude structured events. AMI units are upstream speech segment elements. CaSiNo and AMI counts are source inventories, not processed turns, sampling eligible units, or common analysis denominators.

Under the planned protocol, Dreddit’s official train partition will be the only corpus used to develop prompts, packet rules, controlled variants, routing features, and thresholds. After these choices are frozen, its official test partition will provide an in-domain audit. All 11 CaChe focus groups will be held out from WarrantRoute development for the planned cross-domain confirmatory evaluation. “Held out” refers to the experimental procedure, not model-training provenance. The deposit is public and published themes are known, so we will not call it unopened or model-unseen or use the published themes as scoring gold.

The power analysis and expert time budget determine a common number of base packets for the Dreddit audit and CaChe evaluation. Sampling preserves post clusters in Dreddit and focus group and transcript local speaker clusters in CaChe. Both corpora use identical packet and context limits. Dreddit labels and published CaChe themes may stratify audits but are not treated as thematic ground truth or router features. CaSiNo and AMI enter no generation, sampling, or model selection procedure unless a later preregistration assigns them an experimental role after corpus specific validation and governance approval.

If the governance gates are completed, we will generate packets using [MODEL VERSIONS] and [GENERATION PROCEDURE]. We will retain prompts, decoding settings, context order, run identifiers, source spans, and human revisions and will not select test outputs after inspecting their quality. We will use text or transcripts only, process them under the documented approved governance arrangement, and not send protected material to an external model unless the access terms and our institution permit it. Material released publicly will be synthetic or specifically cleared for release.

[CURRENT INSTITUTIONAL DETERMINATION; SOURCE, PLATFORM, MODEL PROCESSING, AND RATER ACCESS AUTHORIZATIONS FOR DREDDIT, CACHE, CASINO, AND AMI; CASINO LINKAGE RESOLUTION; AMI COMPONENT SPLIT; RATER EXCERPT RULES; FINAL PACKET COUNTS; MODEL SNAPSHOT; PROMPTS; DECODING SETTINGS; RUN COUNTS; AND FROZEN MANIFEST HASHES].

B.2 Role definitions and recruitment

[ELIGIBILITY FOR RESEARCHERS WITHOUT SPECIALIST EXPERTISE, QUALITATIVE METHODS EXPERTS, AND DOMAIN EXPERTS; HANDLING OF DUAL EXPERTISE; RECRUITMENT; COMPENSATION; POSITIONALITY; EXCLUSIONS; AND INDEPENDENCE SAFEGUARDS].

B.3 Complete annotation instrument

The three packet-level quality questions are: (1) *How fully does the cited evidence support the core interpretive meaning in context?* (evidential credibility); (2) *How well does the packet preserve consequential differences, dissent, counterevidence, less-frequent or boundary cases, and contextual qualifications?* (voice-and-boundary preservation); and (3) *How well do the strength and breadth of the claim, including causal or group-level language, match the supplied evidence?* (scope calibration).

For each question, 1 means not adequate, 2 major problems, 3 partially adequate and requiring material revision, 4 mostly adequate with only minor issues, and 5 fully adequate. The tutorial will supply construct-specific examples at each anchor. For rating r_{ijm} from rater i on packet j , we will analyze the full ordinal response; the descriptive adequacy rate for measure m will be the proportion of evaluable rater–packet judgments scored 4 or 5:

Evidential credibility. A 1 indicates absent, unverifiable, misattributed, or contradictory evidence; a 3 indicates that the core meaning is partly supported but a material link or qualifier lacks evidence; and a 5 indicates that every material part of the core interpretation is directly warranted by sufficient contextual evidence.

Voice-and-boundary preservation. A 1 indicates that a consequential voice, counterexample, or contextual distinction is erased

or reversed; a 3 retains the main pattern but underplays a meaningful viewpoint or boundary; and a 5 retains all consequential variation, dissent, less-frequent cases, and contextual qualifications visible in the packet.

Scope calibration. A 1 indicates that claim breadth, polarity, or causal strength is incompatible with the evidence; a 3 has the broadly appropriate direction but requires a material change in scope or strength; and a 5 fully matches participant-, group-, and corpus-level scope and makes no stronger causal or prevalence claim than the evidence permits.

Scores 2 and 4 represent the corresponding intermediate judgments under the common anchors. Raters select “cannot judge” when the supplied excerpts, local context, or source-coverage summary are insufficient for the specified dimension.

$$A_m = \frac{\sum_{i,j} \mathbb{1}(r_{ijm} \in \{4, 5\})}{\sum_{i,j} \mathbb{1}(r_{ijm} \text{ is evaluable})}. \quad (1)$$

“Cannot judge” is neither a midpoint nor a failure and is reported with its own denominator. Confidence has separate anchors from 1 (very uncertain) to 5 (very certain). Exact-string and source-offset checks are recorded as deterministic diagnostics; raters judge whether verified excerpts warrant the interpretation in context.

[CONSENT TEXT, TUTORIAL EXAMPLES, COMPREHENSION CHECKS, ACCEPT/REVISE/REJECT/ESCALATE WORDING, ABSTENTION RULE, RATIONALE PROMPT, AND TIMING DISCLOSURE].

B.4 Artifact-quality definitions and matching

Plural-reference coverage will be calculated only for complete, natural-output code and theme inventories, not from sampled packets or controlled twins. Before inspecting test outputs, multiple analysts will independently produce code and theme concepts. We will freeze their pooled reference space after merging only concepts adjudicated to be equivalent, retaining the originating analyst and source provenance. This pool will represent a bounded set of defensible concepts available to the study team, not universal thematic ground truth.

For output level $\ell \in \{\text{code, theme}\}$, let P_ℓ be the frozen pooled concepts and G_ℓ the generated concepts. Independent experts blinded to system condition may make many-to-many conceptual matches,

denoted by $p \sim g$. Coverage is

$$C_\ell = \frac{|\{p \in P_\ell : \exists g \in G_\ell, p \sim g\}|}{|P_\ell|}. \quad (2)$$

We will report code and theme coverage separately with preregistered uncertainty intervals, list unmatched reference concepts, retain disagreements about matching, and not count unmatched generated concepts as errors without separate evidential review. To diagnose coverage gained through indiscriminate output, we will also report generated code and theme counts, unmatched generated concepts, split and merge patterns, redundancy, and the human-audited unsupported-concept rate. Because the measure relies on the analyst reference space, it will not be available to the router. The three ordinal measures and plural-reference coverage will remain separate; no omnibus quality score will be computed.

For controlled variants, the primary manipulation check will map unsupported evidence to evidential credibility; source concentration, counterevidence loss, and contextual flattening to voice-and-boundary preservation; and unsupported abstraction to scope calibration. We will also report cross-construct effects rather than assume that each defect can affect only one measure. Plural-reference coverage will be restricted to natural-output inventories and will not serve as a manipulation check for controlled twins.

B.5 Packet construction and assignment

[NATURAL ITEM SAMPLING, CONSTRUCTION OF CONTROLLED VARIANTS, VETTING, MANIPULATION CHECKS, RANDOMIZATION, INCOMPLETE BLOCK ASSIGNMENT, AND BLINDING].

B.6 Power and confirmatory analysis

[PILOT VARIANCE COMPONENTS, SIMULATION DESIGN, POWER TARGET, NON-INFERIORITY MARGIN, MODEL EQUATIONS, MULTIPLICITY PLAN, FALLBACKS, AND EXCLUSIONS].

B.7 Routing implementation

[FEATURE DEFINITIONS, PREPROCESSING, MODEL, HYPERPARAMETERS, BUDGET SELECTION, THRESHOLDS, CROSS FITTING, FREEZE DATE, AND TEST ACCESS LOG].

C Detailed Ethics and Data Governance

C.1 Source review and access conditions

Approval of an original collection does not authorize this secondary study, model processing, or disclosure to raters. Before real-text model processing or rater disclosure, we will record our institution’s determination, its scope and date, and the governing source, platform, and access terms for each corpus.

The CaChe source study reports written consent and review by Rush University Medical Center (22011311), Maseno University Ethics Review Committee (MSU/DRPI/MUERC/01021/21), and University of Illinois Chicago (2022-0220), with a favorable opinion from Liverpool School of Tropical Medicine (21-087) (Awiti et al., 2024). The CaSiNo paper reports institutional review, informed consent covering collection and later use, a right to withdraw, and removal of MTurk and HIT identifiers before public release (Chawla et al., 2021). Although that paper calls the released data anonymized, we do not adopt that label: dialogue and rich participant metadata can remain identifying. The AMI site documents European Commission ethics review and signed consent that anticipated possible release to the wider research community; its consent form states that anonymity could not be guaranteed and allowed participants to request removal (AMI Corpus, n.d.). The CC BY 4.0 licenses for CaSiNo and AMI do not grant privacy or personality rights or authorize this project’s model processing, rater disclosure, retention, or release. We did not find an ethics determination, Reddit author consent, privacy procedure, or dataset license in the Dreddit paper or author archive (Turcan and McKeown, 2019); any use therefore remains conditional on a documented institutional and platform basis.

C.2 Project safeguards and rater protections

We do not describe the corpora as anonymous. Dreddit retains searchable verbatim posts. CaChe contains sensitive sexual and reproductive health narratives that may remain identifying despite participant numbers and selective redaction. CaSiNo raw records combine dialogue with demographics, personality measures, preferences, free text reasons, and outcome surveys; those metadata remain excluded from any future working output. AMI source XML contains transcript text, precise timing, and attribute bearing source identifiers; fu-

ture working outputs must omit timings and raw identifiers. Restricted analysis copies use project pseudonymous identifiers, minimize fields, mask direct identifier patterns, and preserve protected linkage information separately. Pattern masking does not replace contextual privacy review. Protected text is not sent to a model or service unless its processing, retention, and training terms are approved. [STATE STORAGE LOCATION, ACCESS ROLES, RETENTION PERIOD, DELETION PROCEDURE, APPROVED MODEL PROVIDERS, AND PROVIDER RETENTION OR TRAINING SETTINGS].

Every excerpt shown to raters or considered for release will receive a two-person privacy and permissions review. Public artifacts will contain aggregate results, code, prompts, schemas, synthetic packets, and only text or ratings specifically cleared for release. We will not publish raw hashes of short sensitive excerpts without assessing linkage risk. Raters will receive a study-specific consent form and privacy warning, may withdraw under the approved procedure, and will participate voluntarily outside supervisory, employment, grading, or authorship decisions. Identifiers will be stored separately from responses. [STATE RECRUITMENT CHANNELS, ELIGIBILITY, COMPENSATION AND RATE JUSTIFICATION, WITHDRAWAL HANDLING, RATER DEMOGRAPHICS, AND FINAL CONSENT LANGUAGE]. Routing is a research review aid, not a clinical or participant-level decision system; random audits and false-acceptance reporting address, but do not remove, the risk of false reassurance.

D Supplementary Results

[ADD CONFUSION MATRICES, SELECTIVE-RISK VERSUS EXPERT-REVIEW-COVERAGE CURVES, RESULTS BY FAILURE FAMILY, COMPARISONS OF NATURAL AND CONTROLLED ITEMS, CONFIDENCE AND ROUTER CALIBRATION PLOTS, RATIONALE ANALYSIS, AND ROBUSTNESS CHECKS].

E Reproducibility and Governance Checklist

- ☐ Confirm institutional, source or platform, model processing, and rater access conditions for Dreddit, CaChe, CaSiNo, and AMI; re-

Evaluator group and construct	Krippendorff α	Gwet AC2	Raw agreement	Cannot judge	Median time
Researchers: credibility	—	—	—	—	—
Researchers: voice/boundary	—	—	—	—	—
Researchers: scope calibration	—	—	—	—	—
Methods experts: credibility	—	—	—	—	—
Methods experts: voice/boundary	—	—	—	—	—
Methods experts: scope calibration	—	—	—	—	—
Domain experts: credibility	—	—	—	—	—
Domain experts: voice/boundary	—	—	—	—	—
Domain experts: scope calibration	—	—	—	—	—

Table 7: Reliability shell for the three ordinal artifact-quality measures. Report intervals and score prevalence in the final appendix; confidence and router calibration are reported separately.

Routing configuration	Dreaddit dev.	Dreaddit audit	CaChE	Recall change	Expert minutes
Full features	—	—	—	—	—
No disagreement	—	—	—	—	—
No abstention	—	—	—	—	—
No provenance	—	—	—	—	—
Confidence only	—	—	—	—	—

Table 8: Effects of transfer and removal of routing features. Define the reported metric in the final caption.

solve CaSiNo linkage or its unsplit role; validate AMI transcript preprocessing and freeze component isolated splits; then freeze the analytic contract, prompts, model snapshots, and manifests.

☐ Obtain ethics determination and finalize voluntary recruitment and withdrawal procedures that avoid coercion.

☐ Freeze role criteria, item manifest, the four metric definitions and anchors, plural-reference construction and matching rules, assignment, exclusions, and power analysis.

☐ Freeze router features, thresholds, budgets measured in expert minutes, primary outcomes, and analysis code before generating CaChE model outputs.

☐ Release only packets cleared by the two-person excerpt review, permitted ratings, prompts, code, model and run metadata, non-text manifest hashes, and text-derived hashes only when source terms permit and linkage-risk review clears them, plus a datasheet.

☐ Replace every `\fillin`, `\blankcell`, and `\blankwide`; set `\predatadraftfalse`; verify the eight page boundary; and update the abstract last.